

TITUS E. C. WORMER

# RETEXT



# RETEXT

Design of an extensible system for analysing and manipulating natural language

TITUS E. C. WORMER

School of Design and Communication  
Communication and Multimedia Design  
Amsterdam University of Applied Sciences

August 2014 – version 0.3

Titus E. C. Wormer: *Retext*, Design of an extensible system for  
analysing and manipulating natural language, © August 2014

SUPERVISOR: Justus Sturkenboom

LOCATION: Delft

---

## ACKNOWLEDGMENTS & DEDICATION

---

*Thanks to my supervisor Justus Sturkenboom for the trust (and the patience) in me and my work, and for allowing me to produce a product I'm pleased with.*

*This thesis is dedicated to the open source community, especially those who raised issues, submitted pull request, and those who haven't done so yet, but will in the future.*



---

## ABSTRACT

---

This document captures the use cases and requirements for designing and standardising a solution for textual manipulation and analysis in ECMAScript. In addition, this paper presents an implementation that meets these requirements.





---

## INTRODUCTION

---

Natural Language Processing (NLP), a field of computer science, artificial intelligence, and linguistics concerned with the interaction between computers and human languages (according to Wikipedia, see ‘Natural language processing’), is becoming more important in society. For example, search engines provide answers before being questioned, the NSA detects threats of violence in text messages, and e-mail applications know if the user forgot to include an attachment.

Despite increased interest, web developers trying to solve NLP problems reinvent the wheel over and over again. There are tools, especially for other platforms—such as in Python (Bird et al.) and Java (Baldrige)—but they either take a too naive approach<sup>1</sup>, or try to do everything out of the box<sup>2</sup>. What is missing is a standard representation of the grammatical hierarchy of text and a standard for multipurpose analysis of natural language.

For over a year I too camped with this problem. On many occasions I tried solving it, to no avail. While working on this thesis, the specification, the product, I developed a well thought out and substantiated solution.

Retext—the implementation introduced in this thesis—and other projects in the Retext family are a new approach to the syntax of natural text. Together they form an extensible system for multipurpose analysis of natural language in ECMAScript.

To achieve this goal, I have organised my paper into five main chapters. In the first chapter I define the scope of this paper and review current implementations, what they lack and where they excel. In the second chapter, I propose a better implementation and show its architectural design. In the third chapter, I define conditions for the production of such a proposal, where I touch upon the target audience, use cases, and requirements.

I end the paper with a fourth chapter which describes the steps taken to validate the proposal. I conclude with a fifth chapter that offers information on expanding the proposal.

Before I can begin the examination of a proposal, I need to provide a context for NLP.

---

<sup>1</sup> Such as ignoring white space (Loadfive), implementing a naive definition of “words” (Hunzaker), or by using an inadequate algorithm to detect sentences (New York Times).

<sup>2</sup> Although a do-all library works well on server-side platforms, it fares less well on the web (such as Umbel et al.), where modularity and moderation are in order.



---

## CONTENTS

---

<b>i</b>	<b>RETEXT</b>	<b>1</b>
<b>1</b>	<b>CONTEXT</b>	<b>3</b>
1.1	Scope . . . . .	4
1.2	Implementations . . . . .	4
1.2.1	Stages . . . . .	4
1.2.2	Tasks . . . . .	5
1.2.3	Using Corpora for NLP . . . . .	7
1.2.4	Using a web API . . . . .	8
<b>2</b>	<b>PRODUCTION</b>	<b>9</b>
2.1	Target Audience . . . . .	9
2.2	Use cases . . . . .	9
2.3	Requirements . . . . .	10
2.3.1	Open Source . . . . .	10
2.3.2	Performance . . . . .	10
2.3.3	Testing . . . . .	11
2.3.4	Code quality . . . . .	11
2.3.5	Automation . . . . .	12
2.3.6	API Design . . . . .	12
2.3.7	Installation . . . . .	13
<b>3</b>	<b>DESIGN &amp; ARCHITECTURE</b>	<b>15</b>
3.1	Syntax: NLCST . . . . .	15
3.2	Parser: Parse-latin . . . . .	16
3.2.1	parse-english . . . . .	17
3.2.2	parse-dutch . . . . .	17
3.3	Object Model: TextOM . . . . .	17
3.4	Natural Language System: Retext . . . . .	18
<b>4</b>	<b>VALIDATION</b>	<b>19</b>
4.1	Plugins . . . . .	19
4.2	Reception . . . . .	19
<b>5</b>	<b>CONCLUSION</b>	<b>21</b>
5.1	Summary . . . . .	21
5.2	Future Work . . . . .	21
5.3	Advice . . . . .	21
<b>ii</b>	<b>APPENDIX</b>	<b>23</b>
<b>A</b>	<b>NLCST DEFINITION</b>	<b>25</b>
<b>B</b>	<b>PARSE-LATIN OUTPUT</b>	<b>29</b>
<b>C</b>	<b>DOM</b>	<b>33</b>
	<b>GLOSSARY</b>	<b>35</b>
	<b>WORKS CITED</b>	<b>37</b>

Part I

RETEXT



---

## CONTEXT

---

{{Needs a nice quote about the definition of NLP}}

The focus of this paper is Natural Language Processing (NLP). NLP is a field related to human–computer interaction, as it concerns itself with enabling machines to understand human language. Human language, a medium which is easy for humans to understand, poses problems for machines.

The Georgetown–IBM experiment in 1945, one of the first application of NLP, illustrates this difficulty. During this study in New York, scientists demonstrated a Russian–English translation system (Hutchins). The machine translated more than sixty sentences from Russian to English. The experiment was well publicised and resulted in optimism. The public believed machine-translation would be a “solved problem” within three to five years. Despite promising initial results, the following ten years proved to be disappointing and led to reduced funding.

Machine translation is just one of many major tasks involved with NLP. Other tasks include generating summaries, detecting references to people and places, or extracting opinion. Tasks which are all part of *information extraction*: the act of finding certain information in a document. Many programs exist to carry out these and many other NLP tasks. The approach taken to perform these tasks are often similar between implementations. Entity linking for example, is often implemented as follows (according to ‘Stanbol’):

1. *Language Detection (optional)* — Based on the language of the given text, the algorithms behind the following steps will change. Omitted if the implementation supports a single language;
2. *Sentence Tokenisation (optional)* — Sentence breaking elevates performance and heightens accuracy of the following stages, in particular pos tagging;
3. *Word Tokenisation* — The entities (words) must be free from their surroundings;
4. *Part-of-Speech (pos) Tagging (optional)* — It is often desired to link several nouns or proper nouns. Detecting word categories makes this achievable;
5. *Noun Phrase Detection (optional)* — Although *apple* and *juice* could be two entities, it is more appropriate to link to one entity: *apple juice*. Detecting noun phrases makes this possible;

6. *Lemmatisation or Stemming (optional)* – Be it *walk*, *walked*, or *walking*, all forms of *walk* could link to the same entity. Detecting either makes this possible;
7. *Entity Linking* – Linking detected entities to references, such as an encyclopaedia.

NLP covers many different tasks, but the process of accomplishing these goals touches, as seen above, on well defined stages.

### 1.1 SCOPE

Although many NLP tasks exist, the standard and the implementation this paper proposes will only cover one: *tokenisation*. Tokenisation, as defined here, includes breaking sentences, words, and other grammatical units.

Another confinement set to scope the proposal, is that it focusses on syntactic grammatical units. Thus, semantic units (i.e., phrases and clauses) are ignored.

In addition, the paper focusses on written language (text), thus ignoring spoken language.

Last, this paper focusses on Latin script languages: written languages using an alphabet based on the the classical Latin alphabet.

### 1.2 IMPLEMENTATIONS

While researching algorithms to tokenise natural language few viable implementations were found. Most algorithms look at either sentence- or word tokenisation (rarely both). This section describes the current implementations, where they excel, and what they lack.

#### 1.2.1 *Stages*

This section delves into how current implementations accomplish tokenisation tasks.

##### 1.2.1.1 *Sentence tokenisation*

Often referred to as sentence boundary disambiguation<sup>1</sup>, sentence tokenisation is an elementary but important part of NLP. It is almost always a stage in NLP applications and not an end goal. Sentence tokenisation makes other stages (e.g., detecting plagiarism or Part-of-Speech (pos) tagging) perform better.

---

<sup>1</sup> Both sentence tokenisation and sentence boundary disambiguation detect sentences. Sentence boundary disambiguation focusses on the position where sentences break (as in, “One sentence?| Two sentences.”, where the pipe symbols refer to the end of one sentence and the beginning of another), whereas sentence tokenisation targets both the start and end location (as in, “{One sentence?} {Two sentences.”, where everything between braces is classified as a sentence).

Oftentimes, sentences end in one of three symbols: either a full stop (.), an interrogative point (?), or an exclamation point (!)<sup>2</sup>. But detecting the boundary of a sentence is not as simple as breaking it at these markers: they might serve other purposes. Full stops often occur in numbers, suffixed to abbreviations or titles, in initialisms<sup>3</sup>, or in embedded content<sup>4</sup>. The interrogative- and exclamation points too can occur ambiguously, such as in a quote (e.g., ‘“Of course!”, she screamed’).

Disambiguation gets even harder when these exceptions *are* in fact a sentence boundary (double negative), such as in “... use the feminine form of idem, ead.” or in ‘“Of course!”, she screamed, “I’ll do it!”’, where in both cases the last terminal marker ends the respective sentence.

#### 1.2.1.2 Word tokenisation

Like sentence tokenisation, word tokenisation is another elementary but important stage in NLP applications. Whether stemming, finding phonetics, or POS tagging, tokenising words is an important precursory step.

Often implementations see words as everything that is *not* white space (i.e., spaces, tabs, feeds) and their boundaries as everything that is (Loadfive).

Some implementations take punctuation marks into account as boundaries. This practice has flaws, as it results in the faulty classification of inter-word punctuation<sup>5</sup> as part of the surrounding word (Umbel et al.).

#### 1.2.2 Tasks

The previous section covered implementations that solve tokenisation stages in NLP applications, such as Natural’s word tokenisers (Umbel et al.). Concluded was that these implementations are lacking. This section covers several implementations that solve these stages as part of a larger task.

- 
- <sup>2</sup> One could argue the in 1962 introduced obscure interrobang (?), used to punctuate rhetorical statements where neither the question nor exclamation alone exactly serve the writer well, should be in this list (Spector).
  - <sup>3</sup> Although the definition of initialism is ambiguous, this paper defines its use as an acronym (an abbreviation formed from initial components, such as “sonar” or “FBI”) with full stops depicting elision (such as “e.g.”, or “K.G.B.”).
  - <sup>4</sup> Embedded content in this paper refers to an external (non-grammatical) value embedded into a grammatical unit, such as a hyperlink or an emoticon. Note that these embedded values often consist of valid words and punctuation marks, but most always shouldn’t be classified as such.
  - <sup>5</sup> Many such inter-word symbols exist, such as hyphenation points, colons (“12:00”), or elision (whether denoted by full stops, “e.g.”; apostrophes, the Dutch “’s”; or slashes, “N/A”).



### 1.2.2.1 *Sentiment Analysis*

Sentiment analysis is an NLP task concerned with the polarity (positive, negative) and subjectivity (objective, subjective) of text. The implementation of sentiment analysis could look as follows:

1. *Detect Language (optional)*;
2. *Sentence Tokenisation (optional)* — Different sentences have different sentiments, tokenising them helps provide better results;
3. *Word Tokenisation* — Needed to compare with the database;
4. *Lemmatisation or Stemming (optional)* — Helps classification;
5. *Sentiment Analysis*.

Sentiment analysers typically include a database mapping either words, stems, or lemmas to their respective polarity and/or subjectivity<sup>6</sup> and return the average sentiment per sentence, or for the whole document. Many implementations exist for this task (Roth; Zimmerman; Sliwinski), many of which do not include inter-word punctuation in their *definition* of words, resulting in less than perfect results<sup>7</sup>.

### 1.2.2.2 *Automatic Summarisation*

Automatic summarisation is an NLP task concerned with the reduction of text to the *major* points retaining the original document. Few open source implementations of automatic summarisation algorithms on the web, in contrast with implementations for sentimental analysis, were found<sup>8</sup>. The implementation of automatic summarisation could look as follows:

1. *Detect Language (optional)*;
2. *Sentence Tokenisation (optional)* — Unless even finer grained control over the document is possible (tokenising phrases), sentences are the smallest unit that should stay intact in the resulting summary;
3. *Word Tokenisation* — Needed to calculate keywords (words which occur more often than expected by chance alone);
4. *Automatic summarisation*.

Automatic summarisers typically return the highest ranking units, be it sentences or phrases, according to several factors:

- A. *Number of words* — An ideal sentence is neither too long nor too short;
- B. *Number of keywords* — Words which occur more often than expected by chance alone in the whole text;

<sup>6</sup> For example, the AFINN database mapping words to polarity (Nielsen).

<sup>7</sup> In fact, all found implementations deploy lacking tokenisations steps. Dubious, as they each create unreachable code through their naivety: all implementations remove dashes from words, while words such as “self-deluded” are included in the databases they use, but never reachable.

<sup>8</sup> For example, on the web only node-summary was found (Brooks), in Scala textteaser was found (Balbin, ‘textteaser’).

- c. *Similarity to title* — Number of words from the document’s title the unit contains;
- d. *Position inside parent* — Initial and final sentences of a paragraph are often more important than sentences buried somewhere in the middle.

Some implementations include only keyword metrics (Brooks), others include all features (Balbin, ‘textteaser’), or even more advanced factors (‘Summly’).

The only implementation working on the web, by James Brook (‘node-summary’), takes a naive sentence tokenisation approach. Such as ignoring sentences terminated by exclamation marks. Both other implementations, and many more, use a whole different approach to sentence tokenisation: Corpora.

### 1.2.3 Using Corpora for NLP

A corpus is a large, structured set of texts used for many NLP and linguistics tasks. Corpora contain items (often words, but sometimes other units) annotated with information (such as POS tags or lemmas).

These colossal (often more than a million words<sup>9</sup>) lumps of data are the basis of many of the newer revolutions in NLP (Mitkov et al.). Parsing based on supervised learning (in NLP, based on annotated corpora), is the opposite of rule-based parsing<sup>10</sup>. Instead of rules (and exceptions to these rules, exceptions to these exceptions, and so on) specified by a developer, supervised learning<sup>11</sup> delegates this task to machines. This delegation results in a more performant, scalable, program.

Parsing based on corpora has proven to be better in several ways over rule-based parsing, but the former has disadvantages:

1. Good training sets are required;
2. If the corpus was created from news articles, algorithms based on it will not fair so well on microblogs (e.g., Twitter posts).
3. Some rule-based approaches for pre- and post processing are still required;

In addition, corpora-based parsing will not work well on the web. Loading corpora over the network each time a user request a web page is unfeasible for most web sites and applications<sup>12</sup>.

<sup>9</sup> The Brown Corpus contains about a million words (Francis and Kučera), the Google N-Gram Corpus contains 155 billion (Brants and Franz).

<sup>10</sup> A simple rule-based sentence tokeniser could be implemented as follows (O’Neil):

- A. If it is a period, it ends a sentence;
- B. If the period is preceded by an abbreviation, it does not end a sentence;
- C. If the next token is capitalised, it ends a sentence.

<sup>11</sup> “[From] a set of labeled examples as training data ... [, make] predictions for all unseen points” (Mohri et al.).

<sup>12</sup> Currently, one technology exists for storing large datasets in a browser: the HTML5 File System API. However, “work on this document has been discontinued”, and the specification “should not be used as a basis for implementation” (Uhrhane).

Two viable alternative approaches exist for the web: rule-based tokenisation, or connecting to a server over the network.

#### 1.2.4 *Using a web API*

Whereas the term Application Programming Interface (API) stands for an interface between two programs, it is often used in web development as requests (from a web browser), and responses (from a web server) over Hypertext Transfer Protocol (HTTP). For example, Twitter has such a service to allow developers to list, replace, create, and delete so-called tweets and other objects (users, images, &c.). This paper uses the term Web API for the latter, and API for any programming interface.

With the rise of the asynchronous web<sup>13</sup>, supervised learning became available through web APIs (Balbin, ‘TextTeaser’; Princeton University; ‘TextRazor’). This made it possible to use supervised learning techniques on the web, without needing to download corpora to a users computer.

However, accessing NLP web APIs over a network has disadvantages. Foremost of which the time involved in sending data over a network and bandwidth used (especially on mobile networks), and heightened security risks.

---

<sup>13</sup> Starting around 2000, Asynchronous JavaScript and XML (AJAX) started to transform the web. Beforehand, significant changes to websites only occurred when a user navigated to a new page. With AJAX however, new content arrived to users without the need for a full page refresh. One of the first examples are the Outlook Web App in 2000 (‘Outlook Web Access - A Catalyst for Web Evolution’) and Gmail in 2004 (Hyder), both examples of how AJAX made the web feel more “app-like”.

---

## PRODUCTION

---

### 2.1 TARGET AUDIENCE

The audience that benefits the most from the proposal, are web developers. Web developers are programmers who specialise in creating software that functions on the world wide web. A group which enables machines to respond to humans. They engage in client side development (building the interface between a human and a machine on the web), and sometimes also in server side development (building the interface between the client side and a server).

Typical areas of work consist of programming in ECMAScript, marking up documents in Hypertext Markup Language (HTML), graphic design through Cascading Style Sheets (css), creating a back end in Node.js, PHP: Hypertext Preprocessor (PHP), or other platforms, contacting a MongoDB, MySQL, or other database, and more.

Additionally, many interdisciplinary skills, such as usability, accessibility, copywriting, information architecture, or optimisation, are also of concern to web developers.

### 2.2 USE CASES

The use cases of the target audience, the web developer, in the field of NLP are many. Research for this paper found several use cases, although it is expected many more could be defined. The tasks below are each categorised into broad, generic fields: analysis, manipulation, and creation.

- A. The developer may intent to summarise natural text (mostly analysis, potentially also manipulation);
- B. The developer may intent to create natural language, e.g., displaying the number of unread messages: “You have 1 unread message,” or “You have 0 unread messages” (creation);
- C. The developer may intent to recognise sentiment in text: is a *tweet* positive, negative, or spam? (analysis);
- D. The developer may intent to replace so-called *dumb* punctuation with *smart* punctuation, such as dumb quotations with (") or (”), three dots with an ellipsis (...), or two hyphens with an en-dash (–) (manipulation);

- E. The developer may intent to count the number of certain grammatical units in a document, such as, words, white space, punctuation, sentences, or paragraphs (analysation);
- F. The developer may intent to recognise the language in which a document is written (analysation);
- G. The developer may intent to find words in a document based on a search term, with regards for the lemma (or stem) and/or phonetics (so that a search for “smit” also returns similar words, such as “Schmidt” or “Smith”) (analysation and manipulation).

NLP is a large field with many challenges, but not every challenge in the field is of interest to the web developer. Foremost, the more academic areas of NLP, such as speech recognition, optical character recognition, text-to-speech transformation, translation, and machine learning, do not fit well with the goals of web developers.

## 2.3 REQUIREMENTS

The proposal must enable the target audience to reach the in the previous section defined use cases. In addition, the proposal should meet several other requirements to better suit the wishes of the target audience.

### 2.3.1 *Open Source*

To reach the target audience and validate its usability, the proposal should be open source. All code should be licensed under MIT License, a license which provides rights for others to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of these projects.

In addition, the software should be developed under the all-seeing eye of the community: GitHub. GitHub is a hosted version control<sup>1</sup> service with social networking features. On GitHub, web developers follow their peers to track what they are working on, watch their favourite projects to get notified of changes, and raise issues and feature requests.

### 2.3.2 *Performance*

The in this paper proposed implementations were produced with high regards for performance. Performance includes the software having a small file size in order to reach the client over the network with the highest possible speed, but most importantly that the execution of code should operate efficiently and at high speeds.

---

<sup>1</sup> Version control services manage revisions to documents, popularly used for controlling and tracking changes in software.

### 2.3.3 *Testing*

With the development of the in this paper introduced software, testing was given high priority. Testing, in software development, refers to validating if software does what it is supposed to do, and can be divided into multiple subgroups:

- A. Unit testing — Validation of each specific section of code;
- B. Integration testing — Validation of how programs work together;
- C. System testing — Validation of if system meets its requirements;
- D. Acceptance testing — Validation of the end product.

Great care was given to develop a full test suite with full coverage for every program. Coverage, in software development, is a term used to describe the amount of code tested by the test suite: full coverage means every part of the code is reached by the tests.

Unit test were run through Mocha (Holowaychuk), coverage was detected with Istanbul (Anantheswaran).

### 2.3.4 *Code quality*

Great attention was given to code quality: how useful and readable for both humans and machines the software is. Special focus was given to consistency and clearness for humans.

#### 2.3.4.1 *Suspicious Code and Bugs*

To detect bugs and suspicious code in the software, ESLint (Zakas) was used. The act of linting, in computer programming, is a term used to describe static code analysis to detect syntactic discrepancies without actually running the code. ESLint was used because it provides a solid basic set of rules and enables developers to create custom rules.

#### 2.3.4.2 *Style*

To enforce a consistent code style, in order to create software readable for humans, JavaScript Code Style Checker (JSCS) was used (Dulin). JSCS provides rules for allowing or disallowing patterns such as white space at the end of a line or camel cased variable names, or setting a maximum line length. JSCS was chosen because it, just like the aforementioned ESLint, provides a strong base set of rules. The rules chosen for the development of the proposed software were set very strict to enforce all code was written in the same manner.

#### 2.3.4.3 *Commenting*

Even when code is completely bug free, uses no hard-to-understand shortcuts, and adheres to a strict style, it might still be hard to understand for humans. The act of commenting code—describing what a program does and why it accomplishes this in a certain way—is

important. On the other hand, commenting can also be too verbose, for example when the code is duplicated in natural language.

JSDoc (‘Annotating JavaScript for the Closure Compiler’) is a markup language for ECMAScript programs, allowing developers to embed documentation in source code. Later various tools can be used to extract the documentation and expose it separately from the original code.

Great care was given to annotate “tricky” source code inside the software with comments, and to apply documentation inside the source code through JSDoc.

### 2.3.5 *Automation*

Great focus was given to develop using several automated Continuous Integration (CI) environments. When suspicious, ambiguous, or buggy code was introduced in the software, the error was automatically detected and in some cases deployment was prevented.

Tools used were Code Climate (‘Code Climate’) to detect complex, duplicate, or bug-prone code, and Travis (‘Travis’) to validate all unit tests passed before deploying the software.

### 2.3.6 *API Design*

Interface design was given high priority for the development of the proposed software. A clear interface of the software, according to Joshua Bloch (‘How to design a good API and why it matters’), has the following characteristics:

1. Easy to learn;
2. Easy to use;
3. Hard to misuse;
4. Easy to read;
5. Easy to maintain;
6. Easy to extend;
7. Meeting its requirements;
8. Appropriate for the target audience.

In essence equal but worded differently are the characteristics of good API design according to the Qt Project (‘API Design Principles’), are as follows:

1. Be minimal;
2. Be complete;
3. Have clear and simple semantics;
4. Be intuitive;
5. Be easy to memorise;
6. Lead to readable code.

With the creation of the software these characteristics, and the in their sources given examples, were taken into account.

### 2.3.7 *Installation*

Access both on the client side and on the server side to the software was of importance during the development of the software. For the server side on Node.js, npm—the default package manager for the platform—is the most popular. On the client side, many different package managers exist, the most popular<sup>2</sup> being Bower and Component. To reach the target audience, in addition of making the whole source available for download through Git and GitHub, NPM, Bower, and Component were used.

---

<sup>2</sup> Popularity here is simply defined as having the most search results on Google.





---

## DESIGN & ARCHITECTURE

---

The in this paper proposed solution to the problem of NLP on the client side is split up in multiple small proposals. Each proposal solves a subproblem.

- A. Natural Language Concrete Syntax Tree (NLCST) — Defines a standard for classifying grammatical units understandable for machines;
- B. Parse-latin — Classifies natural language according to NLCST;
- C. TextOM — Provides an interface for analysing and manipulating output provided by parse-latin;
- D. Retext — Provides an interface for transforming natural language into an object model and exposes an interface for plugins.

The decoupled approach taken by the provided solution enables other developers to provide their own software to replace one of the sub-proposals. For example, other parties can create a parser for the Chinese language and use it instead of parse-latin to classify natural language according to NLCST.

### 3.1 SYNTAX: NLCST

To develop natural language tools in ECMAScript, an intermediate representation of natural language is useful: instead of each module defining their own representation of text, using a single syntax leads to better results, interoperability, and performance.

The elements defined by NLCST (Natural Language Concrete Syntax Tree) are based on the the grammatical hierarchy, but by default do not expose all its constituents<sup>1</sup>. Additionally, more elements are provided to cover other semantic units in natural language<sup>2</sup>.

The definitions of exposed nodes were heavily based on other specifications of syntax trees for manipulation on the web platform, such as CSS, aptly named for the css language (Holowaychuk et al., ‘css’) or the Mozilla JavaScript AST, for ECMAScript (‘Parser API’). Both implementations are widely used. CSS by Rework (Holowaychuk

---

<sup>1</sup> The grammatical hierarchy of text is constituted by words, phrases, clauses, and sentences. NLCSTs only implements the sentence and word constituents by default, although clauses and phrases could be provided by implementations.

<sup>2</sup> Most notably, punctuation, symbol, and white space elements.

et al., ‘rework’), and Mozilla JavaScript Abstract Syntax Tree (AST) by Esprima (Hidayat), Acorn (Haverbeke), and Escodegen (Suzuki).

Note that the aforementioned syntax tree specifications are both ASTs, whereas NLCST is a Concrete Syntax Tree (CST). A concrete syntax tree is a one-to-one mapping of source to result. All information stored in the original input is also available through the resulting tree (Bendersky).

The information stored in CSTs is very verbose and could lead to trees that are hard to work with. On the other hand, the fact that every part of the input is housed in the tree, makes it easy for developers to save the output or pass it on to other libraries for further processing.

See appendix A on page 25 for a complete list of specified nodes of NLCST.

### 3.2 PARSER: PARSE-LATIN

ECMAScript is used extensively. Because of this, many ECMAScript tools are being developed. This includes tools for Natural Language Processing. These ECMAScript tools however, when run on the client-side, can not implement supervised learning based on corpora, and web API usage too is not ideal. Thus, a rule-based parser is needed to tokenise text.

For creating such intermediate representations from Latin-script based languages, parse-latin is presented in this paper<sup>3</sup>. As proof-of-concept, two other libraries are also presented, parse-english and parse-dutch, using parse-latin as a base and providing better support for several language specific features, respectively English and Dutch.

By following NLCSTs, modules building on parse-latin may receive better results or performance over implementing their own parsing tools.

By using the CST as described by NLCST and the parser as described by parse-latin, the intermediate representation can be used by developers to create independent modules.

Basically, parse-latin splits text into white space, word, and punctuation tokens. parse-latin starts out with a pretty simple definition, one that most other tokenisers use:

1. A “word” is one or more letter or number characters;
2. A “white space” is one or more white space characters;
3. A “punctuation” is one or more of anything else.

Then, it manipulates and merges those tokens into a syntax tree, adding sentences and paragraphs where needed.

1. Some punctuation marks are part of the word they occur in, e.g., “non-profit”, “she’s”, “G.I.”, “11:00”, “N/A”;
2. Some full-stops do not mark a sentence end, e.g., “1.”, “e.g.”, “id.”;

<sup>3</sup> Whether Old-English, Icelandic, French, or even scripts slightly similar, such as Cyrillic, Georgian, or Armenian.

3. Although full-stops, question marks, and exclamation marks (sometimes) end a sentence, that end might not occur directly after the mark, e.g., “.)”, “.”.

See appendix B on page 29 for example output provided by the parse-latin parser.

### 3.2.1 *parse-english*

parse-english has the same interface as parse-latin, but returns results better suited for English natural language. For example:

- A. Unit abbreviations (“tsp.”, “tbsp.”, “oz.”, “ft.”, &c.);
- B. Time references (“sec.”, “min.”, “tues.”, “thu.”, “feb.”, &c.);
- C. Business Abbreviations (“Inc.” and “Ltd.”);
- D. Social titles (“Mr.”, “Mmes.”, “Sr.”, &c.);
- E. Rank and academic titles (“Dr.”, “Rep.”, “Gen.”, “Prof.”, “Pres.”, &c.);
- F. Geographical abbreviations (“Ave.”, “Blvd.”, “Ft.”, “Hwy.”, &c.);
- G. American state abbreviations (“Ala.”, “Minn.”, “La.”, “Tex.”, &c.);
- H. Canadian province abbreviations (“Alta.”, “Qué.”, “Yuk.”, &c.);
- I. English county abbreviations (“Beds.”, “Leics.”, “Shrops.”, &c.);
- J. Common elision (omission of letters) (“’n’”, “’o’”, “’em”, “’twas”, “’80s”, &c.).

### 3.2.2 *parse-dutch*

parse-dutch has, just like parse-english, the same interface as parse-latin, but returns results better suited for Dutch natural language. For example:

- A. Unit and time abbreviations (“gr.”, “sec.”, “min.”, “ma.”, “vr.”, “vrij.”, “febr.”, “mrt”, &c.);
- B. Many other common abbreviations: (“Mr.”, “Mv.”, “Sr.”, “Em.”, “bijv.”, “zgn.”, “amb.”, &c.);
- C. Common elision (omission of letters) (“d’”, “’n”, “’ns”, “’t”, “’s”, “’er”, “’em”, “’ie”, &c.).

## 3.3 OBJECT MODEL: TEXTOM

To modify NLCST nodes in ECMAScript, this paper proposes TEXTOM. TEXTOM implements the nodes defined by NLCST, but provides an object-oriented style<sup>4</sup>. TEXTOM was designed to be similar to the Document Object Model (DOM)<sup>5</sup>, the mechanism used by browsers to expose HTML through ECMAScript to developers. Because of TEXTOMs likeness to the DOM, TEXTOM is easy to learn and familiar to the target audience.

<sup>4</sup> Object-oriented programming is a style of programming, where classes, instances, attributes, and methods are important.

<sup>5</sup> See appendix C on page 33 for a more information on the DOM.

TEXTOM provides events (a mechanism for detecting changes), modification functionality (inserting, removing, and replacing children into/from parents), and traversal (e.g., finding all words in a sentence).

NLCST allows authors to extend the specification by defining their own nodes, for example creating phrase or clause nodes. TEXTOM allows for the same extension, and is build to work well with “unknown” node types.

#### 3.4 NATURAL LANGUAGE SYSTEM: RETEXT

For natural language processing on the client side, this paper proposes Retext. Retext combines a parser, such as parse-latin or parse-dutch, with a manipulatable object model (TEXTOM).

In addition, Retext provides a minimalistic plugin mechanism so developers can create and publish plugins for others, and in turn can use others’ plugins inside their projects.

---

## VALIDATION

---

The developed software was validated through two approaches. The design and usability of the interface was validated by trying to solve the use cases with the software. If the target audience actually wanted to use the software was validated through enthusiasm showed by the open source community.

### 4.1 PLUGINS

More than fifteen (15) Retext plugins were created to validate how the different projects integrated together, and how the system worked as a whole.

The created plugins include tools for:

- A. Transforming so-called dumb punctuation marks into more typographically correct punctuation marks ('retext-smartypants');
- B. Transforming emoji short-codes (:cat:) into real emoji ('retext-emoji');
- C. Detecting the direction of text ('retext-directionality');
- D. Detecting phonetics ('retext-double-metaphone');
- E. Detecting the stem of words ('retext-porter-stemmer');
- F. Detecting grammatical units ('retext-visit');
- G. Finding text, even misspelled ('retext-search');
- H. Detecting POS tags ('retext-pos');
- I. Finding keywords and -phrases ('retext-keywords').

The creation of these plugins brought several problems to light in the developed software. These problems were then dealt with back and forth between the software and the plugins. The software changed severely by these fixed problems, which resulted in a more useable interface.

### 4.2 RECEPTION

To validate if the target audience actually wanted to use the developed software, several blogs and email newsletters were contacted to feature Retext, either in the form of an article or as a simple link.

This resulted in several mentions on blogs (Young), link roundups (e.g., Misiti; Sorhus), reddit (e.g., Polencic; Wormer, 'Natural Language Parsing with Retext'; 'DailyJS: Natural Language Parsing with Retext').

In turn, these publications resulted in positive reactions, such as on Twitter (dailyjs; Ahmed; Oswald; Rinaldi; Grigorik), feedback (gut4; Gonzaga dos Santos Filho), and fixes (Burkhead). In addition, many web developers started following the project on GitHub ('Stargazers').

# 5

---

## CONCLUSION

---

### 5.1 SUMMARY

### 5.2 FUTURE WORK

### 5.3 ADVICE





Part II

APPENDIX





---

## NLCST DEFINITION

---

### NODE

Node represents any actual unit in the NLCST hierarchy.

```
interface Node {  
    type: string;  
}
```

### PARENT

Parent (Node) represents a unit in the NLCST hierarchy which can have zero or more children.

```
interface Parent <: Node {  
    children: [];  
}
```

### TEXT

Text (Node) represents a unit in the NLCST hierarchy which has a value.

```
interface Text <: Node {  
    value: string | null;  
    location: Location | null;  
}
```

### LOCATION

Location represents the node's actual location in the source input.

```
interface Location {  
    start: Position;  
    end: Position;  
}
```

### POSITION

Position represents the a position in the source input.

```
interface Position {
    line: uint32 >= 1;
    column: uint32 >= 1;
}
```

#### ROOTNODE

Root (Parent) represents the document.

```
interface RootNode < Parent {
    type: "RootNode";
}
```

#### PARAGRAPHNODE

Paragraph (Parent) represents a self-contained unit of a discourse in writing dealing with a particular point or idea.

```
interface ParagraphNode < Parent {
    type: "ParagraphNode";
}
```

#### SENTENCENODE

Sentence (Parent) represents a grouping of grammatically linked words, that in principle tells a complete thought (although it may make little sense taken in isolation out of context).

```
interface SentenceNode < Parent {
    type: "SentenceNode";
}
```

#### WORDNODE

Word (Parent) represents the smallest element that may be uttered in isolation with semantic or pragmatic content.

```
interface WordNode < Parent {
    type: "WordNode";
}
```

#### PUNCTUATIONNODE

Punctuation (Parent) represents typographical devices which aid the understanding and correct reading of other grammatical units.

```
interface PunctuationNode < Parent {
    type: "PunctuationNode";
}
```

## WHITESPACENODE

White Space (Punctuation) represents typographical devices devoid of content, separating other grammatical units.

```
interface WhiteSpaceNode < PunctuationNode {
    type: "WhiteSpaceNode";
}
```

## SOURCENODE

Source (Text) represents an external (non-grammatical) value embedded into a grammatical unit, for example a hyperlink or an emoticon.

```
interface SourceNode < Text {
    type: "SourceNode";
}
```

## TEXTNODE

Text (Text) represents actual content in a NLAST document: One or more characters.

```
interface TextNode < Text {
    type: "TextNode";
}
```



# B

---

## PARSE-LATIN OUTPUT

---

An example of how parse-latin tokenises the paragraph “A simple sentence. Another sentence.”, is represented as follows,

```
{
  "type": "RootNode",
  "children": [
    {
      "type": "ParagraphNode",
      "children": [
        {
          "type": "SentenceNode",
          "children": [
            {
              "type": "WordNode",
              "children": [{
                "type": "TextNode",
                "value": "A"
              }]
            },
            {
              "type": "WhiteSpaceNode",
              "children": [{
                "type": "TextNode",
                "value": " "
              }]
            },
            {
              "type": "WordNode",
              "children": [{
                "type": "TextNode",
                "value": "simple"
              }]
            },
            {
              "type": "WhiteSpaceNode",
              "children": [{
                "type": "TextNode",
                "value": " "
              }]
            }
          ]
        }
      ]
    }
  ]
}
```



```

    },
    {
      "type": "WordNode",
      "children": [{
        "type": "TextNode",
        "value": "sentence"
      }]
    },
    {
      "type": "PunctuationNode",
      "children": [{
        "type": "TextNode",
        "value": "."
      }]
    }
  ]
},
{
  "type": "WhiteSpaceNode",
  "children": [
    {
      "type": "TextNode",
      "value": " "
    }
  ]
},
{
  "type": "SentenceNode",
  "children": [
    {
      "type": "WordNode",
      "children": [{
        "type": "TextNode",
        "value": "Another"
      }]
    },
    {
      "type": "WhiteSpaceNode",
      "children": [{
        "type": "TextNode",
        "value": " "
      }]
    },
    {
      "type": "WordNode",
      "children": [{
        "type": "TextNode",
        "value": "sentence"
      }]
    }
  ]
}

```

```

    },
    {
      "type": "PunctuationNode",
      "children": [{
        "type": "TextNode",
        "value": "."
      }]
    }
  ]
}

```





---

## DOM

---

The DOM specification defines a platform-neutral model for errors, events, and (for this paper, the primary feature) node trees. XML-based documents can be represented by the DOM.

Consider the following HTML document:

```
<!DOCTYPE html>
<html class=e>
  <head><title>Aliens?</title></head>
  <body>Why yes.</body>
</html>
```

Is represented by the DOM as follows:

```
|– Document
  |– Doctype: html
  |– Element: html class="e"
    |– Element: head
      |– Element: title
        |– Text: Aliens?
    |– Text: \ n
    |– Element: body
      |– Text: Why yes.\n
```

The DOM interfaces of bygone times were widely considered horrible, but newer features seem to be gaining popularity in the web authoring community as broader implementation across user agents is reached.



---

## GLOSSARY

---

AJAX	Asynchronous JavaScript and XML. 8
API	Application Programming Interface. 8, 12, 16
AST	Abstract Syntax Tree. 15, 16
CI	Continuous Integration. 12
CSS	Cascading Style Sheets. 9, 15
CST	Concrete Syntax Tree. 16
DOM	Document Object Model. 17, 33
ECMAScript	More commonly known as JavaScript (which is in fact a proprietary eponym), ECMAScript is a language widely used for client-side programming on the web. vii, ix, 9, 12, 15–17
HTML	Hypertext Markup Language. 9, 17
HTML5	Hypertext Markup Language, version 5. 8
HTTP	Hypertext Transfer Protocol. 8
IBM	International Business Machines Corporation, a U.S. multinational technology and consulting corporation. 3
JSCS	JavaScript Code Style Checker. 11
JSDoc	JavaScript Documentation, markup language using comments to annotate ECMAScript source code. 12
MIT	Michigan Institute of Technology. 10
MongoDB	Mongo Database. 9
MySQL	My Structured Query Language. 9
NLCST	Natural Language Concrete Syntax Tree. 15–18, 25
NLP	Natural Language Processing. ix, 3–10, 15
npm	Package manager for, and included in, Node.js. 13
NSA	National Security Agency, a U.S. intelligence agency. ix
PHP	PHP: Hypertext Preprocessor. 9

POS	Part-of-Speech. 3–5, 7, 19
TextOM	Text Object Model. 15, 17, 18

---

## WORKS CITED

---

- Ahmed, S. [sarfraznawaz]. ‘New #dailyjs #javascript post : Natural Language Parsing with Retext <http://ift.tt/1qrUjGo>’.  
Twitter, 31/7/2014. Web. 8th Aug. 2014.
- Anantheswaran, K. [gotwarlost]. ‘istanbul’. *gotwarlost/istanbul*.  
GitHub, 5/8/2014, version 0.3.0. Web. 9th Aug. 2014.
- ‘Annotating JavaScript for the Closure Compiler’.  
*Google Developers: Closure Tools*. Google, 30/6/2014.  
Web. 8th Aug. 2014.
- ‘API Design Principles’. *Qt Project: Qt Wiki*.  
Qt Project, Digia Plc, 7/8/2014. Web. 8th Aug. 2014.
- Balbin, J. ‘TextTeaser’.  
*TextTeaser: An Automatic Summarization Application and API*.  
TextTeaser, 2014. Web. 8th Aug. 2014.
- [Mojojolo]. ‘textteaser’. *Mojojolo/textteaser*. GitHub, 25/6/2014.  
Web. 9th Aug. 2014.
- Baldrige, J. ‘OpenNLP’.  
*Apache OpenNLP: Welcome to Apache OpenNLP!* Apache, 2005.  
Web. 9th Aug. 2014.
- Bendersky, E. ‘Abstract vs. Concrete Syntax Trees’.  
*Eli Bendersky’s Website*. 16/2/2009. Web. 8th Aug. 2014.
- Bird, S., E. Klein and E. Loper.  
*Natural Language Processing with Python*. 1st ed.  
Sebastopol, California: O’Reilly Media, 7/2009. Print.
- Bloch, J. ‘How to design a good API and why it matters’.  
*Companion to the 21st ACM SIGPLAN symposium on Object-oriented programming systems, languages, and applications*. ACM. 2006.  
506–507. Print.
- Brants, T. and A. Franz. ‘{Web 1T 5-gram Version 1}’ (2006). Print.
- Brooks, J. [jbrooksuk]. ‘node-summary’. *jbrooksuk/node-summary*.  
GitHub, 18/2/2014, version 1.0.0. Web. 9th Aug. 2014.
- Burkhead, J. [jlburkhead]. ‘Spelling fix in README’.  
*wooorm/retext: Pull Request #11*. GitHub, 4/8/2014.  
Web. 8th Aug. 2014.
- ‘Code Climate’. *Code Climate: Hosted static analysis for Ruby and JavaScript source code*. Bluebox, web. 8th Aug. 2014.
- dailyjs [dailyjs]. ‘Natural Language Parsing with Retext: <http://dailyjs.com/2014/07/31/retext>’. Twitter, 31/7/2014.  
Web. 8th Aug. 2014.
- Dulin, M. [mdevils]. ‘JSCS’. *mdevils/node-jscs*.  
GitHub, 8/8/2014, version 1.5.9. Web. 9th Aug. 2014.
- Francis, W. N. and H. Kučera. ‘Brown corpus manual’.  
*Brown University Department of Linguistics* (1979). Print.



- Gonzaga dos Santos Filho, L. [Ifilho]. 'Is this English only?'  
*woorm/retext: Issue #14*. GitHub, 8/8/2014. Web. 9th Aug. 2014.
- Grigorik, I. [igrigorik].  
 'English (latin) language parser in JavaScript: <http://bit.ly/V8cCq7> - aka, English > AST > transform and ... > profit. fun stuff!'  
 Twitter, 8/8/2014. Web. 9th Aug. 2014.
- gut4. 'Some demos not working in chrome >36'.  
*woorm/retext: Issue #10*. GitHub, 1/8/2014. Web. 8th Aug. 2014.
- Haverbeke, M. [marijnh]. 'acorn'. *marijnh/acorn*.  
 GitHub, 31/7/2014, version 0.6.0. Web. 9th Aug. 2014.
- Hidayat, A. [ariya]. 'esprima'. *ariya/esprima*.  
 GitHub, 30/7/2014, version 1.2.2. Web. 9th Aug. 2014.
- Holowaychuk, T. [visionmedia]. 'mocha'. *visionmedia/mocha*.  
 GitHub, 6/8/2014, version 1.2.1.4. Web. 9th Aug. 2014.
- Holowaychuk, T., S. Baumgartner, J. Ong, K. Mårtensson, M. Bu, M. Thirouin, N. Gallagher, A. Sexton and (dead-horse) [reworkcss].  
 'css'. *reworkcss/css*. GitHub, 5/8/2014, version 2.1.0.  
 Web. 9th Aug. 2014.
- [reworkcss]. 'rework'. *reworkcss/rework*.  
 GitHub, 25/6/2014, version 1.0.0. Web. 9th Aug. 2014.
- Hunzaker, N. [nhunzaker]. 'speakeasy'. *nhunzaker/speakeasy*.  
 GitHub, 10/4/2013, version 0.2.11. Web. 9th Aug. 2014.
- Hutchins, J. 'The first public demonstration of machine translation: the Georgetown-IBM system, 7th January 1954'. Expanded version of the paper presented at the AMTA conference in September 2004. Print.
- Hyder, Z. 'Gmail: 9 Years and Counting'. *Official Gmail Blog*.  
 Google, 10/4/2013. Web. 8th Aug. 2014.
- Loadfive [loadfive]. 'Knwl.js'. *loadfive/knwl.js*.  
 GitHub, 4/8/2014, version 0.0.1. Web. 9th Aug. 2014.
- Misiti, J. [josephmisiti]. 'Awesome Machine Learning'.  
 GitHub, 3/8/2014. Web. 8th Aug. 2014.
- Mitkov, R., C. Orasan and R. Evans.  
 'The importance of annotated corpora for NLP: the cases of anaphora resolution and clause splitting'. *Proceedings of Corpora and NLP: Reflecting on Methodology Workshop*. Citeseer.  
 Cargese, Corse, 7/1999. 60–69. Print.
- Mohri, M., A. Rostamizadeh and A. Talwalkar. *Foundations of Machine Learning (Adaptive Computation and Machine Learning series)*.  
 MIT press, 8/2012. Print.
- 'Natural language processing'. *Wikipedia: The Free Encyclopedia*.  
 Wikimedia Foundation, Inc., 5/8/2014. Web. 8th Aug. 2014.
- New York Times [NYTimes]. 'Emphasis'. *NYTimes/Emphasis*.  
 GitHub, 5/4/2012. Web. 9th Aug. 2014.
- Nielsen, F. Å. 'AFINN' (3/2011). Print.
- O'Neil, J. 'Doing Things with Words, Part Two: Sentence Boundary Detection'. *Attivo*. Attivo, 29/10/2008. Web. 8th Aug. 2014.

- Oswald, T. [therebelrobot]. 'This. Just. Made. My. Brain. Explode.  
<https://github.com/wooorm/retext>'. Twitter, 1/8/2014.  
 Web. 8th Aug. 2014.
- 'Outlook Web Access - A Catalyst for Web Evolution'.  
*The Exchange Team Blog: Blogs*. Microsoft Corp., 21/6/2005.  
 Web. 8th Aug. 2014.
- 'Parser API'. *Mozilla: MDN*. Mozilla, 29/7/2014. Web. 8th Aug. 2014.
- Polencic, D. [danielepolencic]. 'Extensible system for analysing and  
 manipulating natural language'. *Node*. Reddit, 29/7/2014.  
 Web. 8th Aug. 2014.
- Princeton University. 'About WordNet'. *WordNet: About WordNet*.  
 Princeton University, 2010. Web. 8th Aug. 2014.
- Rinaldi, B. [remotesynth]. 'Retext is a JavaScript natural language  
 parser useful for doing things like removing profanity, analyzing  
 text, etc. <http://bit.ly/1zCZ1GP>'. Twitter, 1/8/2014.  
 Web. 8th Aug. 2014.
- Roth, K. [thinkroth]. 'Sentimental'. *thinkroth/Sentimental*.  
 GitHub, 12/4/2014, version 1.0.1. Web. 9th Aug. 2014.
- Sliwinski, A. [thisandagain]. 'sentiment'. *thisandagain/sentiment*.  
 GitHub, 7/6/2014, version 1.0.1. Web. 9th Aug. 2014.
- Sorhus, S. [sindresorhus]. 'Awesome Node.js'. GitHub, 7/8/2014.  
 Web. 8th Aug. 2014.
- Spector, P. 'Welcome to Interrobang-Mks'.  
*Welcome to Interrobang-Mks*.  
 American Translators Association, 1/1/1996. Web. 8th Aug. 2014.
- 'Stanbol'. *Apache Stanbol: Welcome to Apache Stanbol!*  
 Apache, 2010-12-14. Web. 9th Aug. 2014.
- 'Stargazers'. *wooorm/retext*. GitHub, 8/8/2014. Web. 8th Aug. 2014.
- 'Summly'. *Summly: Pocket sized news for iPhone*. Summly, 2012.  
 Web. 8th Aug. 2014.
- Suzuki, Y. [constellation]. 'escodegen'. *constellation/escodegen*.  
 GitHub, 28/7/2014, version 1.3.3. Web. 9th Aug. 2014.
- 'TextRazor'. *TextRazor: The Natural Language Processing API*.  
 TextRazor, 2014. Web. 8th Aug. 2014.
- 'Travis'. *Travis*. Travis CI GmbH, 2014. Web. 8th Aug. 2014.
- Umbel, C., R. Ellis and K. Koch. 'natural'. *NaturalNode/natural*.  
 GitHub, 22/6/2014, version 0.1.28. Web. 9th Aug. 2014.
- Wormer, T. [wooorm\_].  
 'DailyJS: Natural Language Parsing with Retext'. *JavaScript*.  
 Reddit, 5/8/2014. Web. 8th Aug. 2014.
- [wooorm\_]. 'Natural Language Parsing with Retext'. *Node*.  
 Reddit, 1/8/2014. Web. 8th Aug. 2014.
- [wooorm]. 'retext-directionality'. *wooorm/retext-directionality*.  
 GitHub, 22/7/2014, version 0.1.1. Web. 9th Aug. 2014.
- [wooorm]. 'retext-double-metaphone'.  
*wooorm/retext-double-metaphone*. GitHub, 13/7/2014, version 0.1.0.  
 Web. 9th Aug. 2014.

- [wooorm]. 'retext-emoji'. *wooorm/retext-emoji*.  
GitHub, 13/7/2014, version 0.1.0. Web. 9th Aug. 2014.
- [wooorm]. 'retext-keywords'. *wooorm/retext-keywords*.  
GitHub, 16/7/2014, version 0.0.1. Web. 9th Aug. 2014.
- [wooorm]. 'retext-porter-stemmer'.  
*wooorm/retext-porter-stemmer*. GitHub, 13/7/2014, version 0.1.0.  
Web. 9th Aug. 2014.
- [wooorm]. 'retext-pos'. *wooorm/retext-pos*.  
GitHub, 31/7/2014, version 0.1.2. Web. 9th Aug. 2014.
- [wooorm]. 'retext-search'. *wooorm/retext-search*.  
GitHub, 13/7/2014, version 0.1.0. Web. 9th Aug. 2014.
- [wooorm]. 'retext-smartypants'. *wooorm/retext-smartypants*.  
GitHub, 13/7/2014, version 0.1.0. Web. 9th Aug. 2014.
- [wooorm]. 'retext-visit'. *wooorm/retext-visit*.  
GitHub, 7/7/2014, version 0.1.0. Web. 9th Aug. 2014.
- Young, A. 'Natural Language Parsing with Retext'.  
*DailyJS: A JavaScript Blog*. 31/7/2014. Web. 8th Aug. 2014.
- Zakas, N. [eslint]. 'eslint'. *eslint/eslint*. GitHub, 6/8/2014, version 0.7.4.  
Web. 9th Aug. 2014.
- Zimmerman, M. [mileszim]. 'sediment'. *mileszim/sediment*.  
GitHub, 9/9/2013, version 0.9.2. Web. 9th Aug. 2014.